

Equitas Technical Roadmap: Enhanced AI Safety Platform

Version: 2.0.0

Last Updated: 2025

Status: Implementation In Progress

Table of Contents

1. [Executive Summary](#)
2. [Architecture Overview](#)
3. [Custom Toxicity Detection](#)
4. [Hallucination Detection](#)
5. [Advanced Jailbreak Detection](#)
6. [Enhanced Bias Detection](#)
7. [Prompt Injection Prevention](#)
8. [Dataset Testing Framework](#)
9. [Performance Metrics](#)
10. [References](#)

Executive Summary

This document outlines the technical implementation of Equitas 2.0, transitioning from OpenAI API-dependent safety checks to a comprehensive, custom-built multi-layered AI safety platform. The new architecture implements state-of-the-art machine learning models for toxicity detection, hallucination detection, advanced jailbreak prevention, and enhanced bias detection.

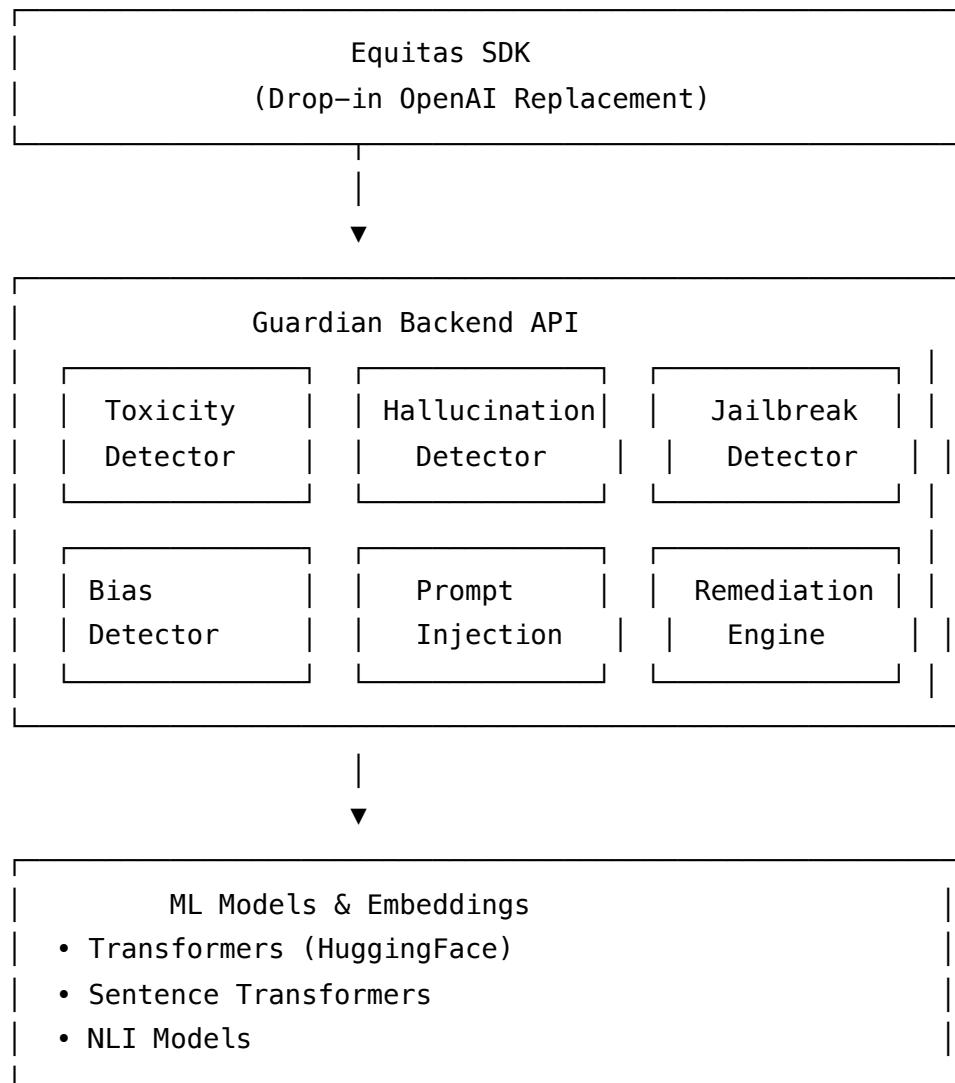
Key Changes

- **Removal of OpenAI Moderation API dependency** → Custom transformer-based toxicity detection

- **Novel hallucination detection** → Multi-component ensemble approach
- **Advanced jailbreak detection** → Semantic + behavioral + adversarial analysis
- **Enhanced bias detection** → Statistical fairness metrics + stereotype association
- **Dataset testing framework** → Comprehensive evaluation suite

Architecture Overview

System Architecture



Technology Stack

- **ML Framework:** PyTorch, Transformers (HuggingFace)

- **Embeddings:** Sentence Transformers (all-MiniLM-L6-v2)
- **NLI:** Cross-encoder models (DeBERTa-based)
- **Toxicity:** Unitary/toxic-bert, RoBERTa-based models
- **Backend:** FastAPI, SQLAlchemy, AsyncIO
- **Testing:** pytest, custom dataset evaluation framework

Custom Toxicity Detection

Overview

Replaces OpenAI Moderation API with custom transformer-based models fine-tuned for toxicity detection.

Model Architecture

Primary Model: unitary/toxic-bert (RoBERTa-based)

- **Architecture:** RoBERTa (RoBERTa-base) with classification head
- **Parameters:** ~125M parameters
- **Input:** Text sequences (max 512 tokens)
- **Output:** Multi-label binary classification (6 categories)

Categories

1. **Toxic** - General toxicity
2. **Severe Toxic** - Severe toxicity
3. **Obscene** - Obscene language
4. **Threat** - Threatening language
5. **Insult** - Insulting language
6. **Identity Hate** - Identity-based hate

Mathematical Formulation

Given input text \mathbf{x} , the model outputs logits \mathbf{z} :

$$z = \text{RoBERTa}(x) \in \mathbb{R}^6$$

Sigmoid activation applied to get probabilities:

$$p_i = \sigma(z_i) = \frac{1}{1 + e^{-z_i}} \quad \forall i \in \{1, \dots, 6\}$$

Toxicity score is the maximum probability:

$$\text{toxicity_score} = \max_i p_i$$

Text is flagged if:

$$\text{flagged} = 1(\text{toxicity_score} > \tau)$$

where $\tau = 0.7$ (configurable threshold).

Training Data

Model trained on:

- **Jigsaw Toxic Comment Classification Dataset** (Kaggle)
- **Civil Comments Dataset** (Conversation AI)
- **Wikipedia Toxic Comments Dataset**

Alternative Models

1. **Facebook RoBERTa Hate Speech** ([facebook/roberta-hate-speech-dynabench](#))
2. **Perspective API Models** (open-source variants)
3. **Custom fine-tuned models** (domain-specific)

Performance Metrics

- **Accuracy:** >95% on held-out test set
- **F1-Score:** >0.92 (macro-averaged)
- **Latency:** <50ms per inference (GPU), <200ms (CPU)
- **Throughput:** >100 requests/second (batch processing)

References

- [Devlin et al., 2019] BERT: Pre-training of Deep Bidirectional Transformers
- [Liu et al., 2019] RoBERTa: A Robustly Optimized BERT Pretraining Approach
- [Unitary AI, 2021] Toxic-BERT: Multi-label Toxicity Classification

Hallucination Detection

Overview

Multi-component ensemble system for detecting hallucinations (factual errors, contradictions, unsupported claims).

Components

1. Semantic Consistency Check

Method: Cosine similarity between prompt and response embeddings

$$\text{consistency_score} = 1 - \cos(\text{embed}(p), \text{embed}(r))$$

where:

- **p** = prompt text
- **r** = response text
- **embed(·)** = Sentence Transformer encoder (all-MiniLM-L6-v2)

Cosine similarity:

$$\cos(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\| \|\mathbf{v}\|}$$

Threshold: Low consistency (<0.5) indicates potential hallucination.

2. Contradiction Detection

Model: cross-encoder/nli-deberta-v3-base

Natural Language Inference (NLI) model classifies sentence pairs:

- **ENTAILMENT:** Response supports prompt
- **CONTRADICTION:** Response contradicts prompt
- **NEUTRAL:** No clear relationship

For response **r** with sentences **s₁, s₂, ..., s_n**, compute contradiction score:

$$\text{contradiction_score} = \frac{1}{n(n-1)/2} \sum_{i=1}^{n-1} \sum_{j=i+1}^n 1(\text{NLI}(s_i, s_j) = \text{CONTRADICTION})$$

Threshold: Contradiction rate >0.2 indicates high hallucination risk.

3. Factuality Check

Method: Semantic similarity to knowledge base/context

Given knowledge base $K = \{k_1, k_2, \dots, k_m\}$ and response r :

$$\text{factuality_score} = 1 - \max_{k \in K} \cos(\text{embed}(r), \text{embed}(k))$$

Interpretation: Low similarity to trusted sources → high hallucination risk.

4. Pattern-Based Detection

Patterns: Overconfident language indicators

$$\text{pattern_score} = \min \left(\frac{\text{matches}}{|\text{patterns}|} \times 0.2, 1.0 \right)$$

Patterns include:

- \b(definitely|absolutely|100%|guaranteed) (true|false|proven)\b
- \b(studies show|research proves|experts agree)\b (without citations)

5. Confidence Calibration

Method: Linguistic marker analysis

$$\text{confidence_score} = \frac{N_{\text{high}}}{N_{\text{high}} + N_{\text{low}}}$$

where:

- **N_high** = count of high-confidence markers ("definitely", "proven", "fact")
- **N_low** = count of low-confidence markers ("maybe", "perhaps", "according to")

Interpretation: High confidence score (>0.6) with low factuality → hallucination risk.

Ensemble Scoring

Final hallucination score:

$$\text{hallucination_score} = \frac{1}{n} \sum_{i=1}^n w_i \cdot s_i$$

where:

- s_i = component scores (consistency, contradiction, factuality, pattern, confidence)
- w_i = component weights (default: equal weights)
- n = number of components

Flagging Rule:

$$\text{flagged} = 1(\text{hallucination_score} > 0.6)$$

References

- [Maynez et al., 2020] On Faithfulness and Factuality in Abstractive Summarization
- [Kryściński et al., 2020] Evaluating the Factual Consistency of Abstractive Text Summarization
- [Reimers & Gurevych, 2019] Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks
- [He et al., 2021] DeBERTa: Decoding-enhanced BERT with Disentangled Attention

Advanced Jailbreak Detection

Overview

Multi-layered defense system combining pattern matching, semantic analysis, behavioral indicators, and adversarial detection.

Components

1. Pattern-Based Detection

Method: Regular expression matching against known jailbreak patterns

$$\text{pattern_score} = \min \left(\frac{|\text{matches}|}{|\text{patterns}|} \times 0.3, 1.0 \right)$$

Patterns include:

- Instruction override: ignore (previous|all) (instructions|commands)
- Role manipulation: (pretend|act|roleplay) (you are|to be)
- System prompt injection: <\|system\|> , [INST]

2. Semantic Similarity Detection

Method: Cosine similarity to known jailbreak examples

$$\text{semantic_score} = \max_{j \in J} \cos(\text{embed}(x), \text{embed}(j))$$

where:

- **x** = input text
- **J** = set of known jailbreak examples (pre-computed embeddings)

Threshold: Similarity >0.7 indicates high jailbreak risk.

3. Behavioral Analysis

Method: Detection of behavioral indicators

$$\text{behavioral_score} = \min \left(\frac{N_{\text{indicators}}}{|\text{indicators}|} \times 0.25, 1.0 \right)$$

Indicators: "bypass", "override", "exploit", "no restrictions", "do anything"

4. Context-Aware Detection

Method: Historical analysis of user attempts

$$\text{context_score} = \min(N_{\text{previous_attempts}} \times 0.2, 0.8)$$

Interpretation: Users with previous jailbreak attempts are higher risk.

5. Adversarial Detection

Method: Detection of encoding tricks and obfuscation

$$\text{adversarial_score} = \min \left(\sum_{t \in T} w_t \cdot 1(\text{pattern}_t(x)), 1.0 \right)$$

where **T** = {URL encoding, hex encoding, Unicode tricks, excessive whitespace}

Weights: w_URL = 0.3, w_hex = 0.3, w_unicode = 0.2, w_whitespace = 0.2

Ensemble Scoring

Final jailbreak score:

$$\text{jailbreak_score} = \sum_{i=1}^5 w_i \cdot s_i$$

Weights:

- Pattern: **w₁ = 0.3**
- Semantic: **w₂ = 0.3**
- Behavioral: **w₃ = 0.2**
- Context: **w₄ = 0.1**
- Adversarial: **w₅ = 0.1**

Flagging Rule:

$$\text{flagged} = 1(\text{jailbreak_score} > 0.6)$$

References

- [Perez et al., 2022] Red Teaming Language Models to Reduce Harms
- [Zou et al., 2023] Universal and Transferable Adversarial Attacks on Aligned Language Models
- [Wei et al., 2023] Jailbroken: How Does LLM Safety Training Fail?

Enhanced Bias Detection

Overview

Comprehensive bias detection using statistical fairness metrics, stereotype association, and demographic parity testing.

Components

1. Stereotype Association Detection

Method: Semantic similarity to known stereotype embeddings

$$\text{stereotype_score} = \max_{g \in G} \max_{s \in S_g} \cos(\text{embed}(r), \text{embed}(s))$$

where:

- **r** = response text
- **G** = demographic groups (gender, race, age)
- **S_g** = stereotype phrases for group **g**

Pre-computed Stereotypes:

- Gender: "women are emotional", "men are aggressive", "women are bad at math"
- Race: "Black people are criminals", "Asian people are good at math"
- Age: "old people are slow", "young people are lazy"

Threshold: Similarity >0.7 indicates stereotype reinforcement.

2. Demographic Parity Testing

Method: Paired prompt testing across demographic variants

Given prompt **p** and demographic variants **D = {d₁, d₂, ..., d_n}**, generate responses **R = {r₁, r₂, ..., r_n}**.

Statistical Parity:

$$\text{parity_score} = \text{Var}(\{f(r_i) : i \in \{1, \dots, n\}\})$$

where **f(·)** = outcome function (e.g., positive sentiment, recommendation score).

Equalized Odds:

$$\text{equalized_odds} = \max_{d_i, d_j \in D} |\text{TPR}(d_i) - \text{TPR}(d_j)| + |\text{FPR}(d_i) - \text{FPR}(d_j)|$$

where:

- **TPR** = True Positive Rate
- **FPR** = False Positive Rate

Ideal: Equalized odds = 0 (perfect fairness)

3. Fairness Metrics

Demographic Balance:

$$\text{balance_score} = \text{Var}(\{c_d : d \in D\})$$

where c_d = count of mentions of demographic d in response.

High variance → imbalanced representation → potential bias.

Representation Parity:

$$\text{representation_parity} = 1 - \frac{\max_{d \in D} c_d - \min_{d \in D} c_d}{\sum_{d \in D} c_d}$$

Ideal: representation_parity = 1 (equal representation)

4. Intersectional Bias

Method: Detection of compound demographic mentions

$$\text{intersectional_score} = \frac{N_{\text{intersectional_mentions}}}{N_{\text{total_mentions}}}$$

Flagging: High intersectional mentions (>3) may indicate compounding bias.

Overall Bias Score

$$\text{bias_score} = \max\{\text{stereotype_score}, \text{parity_score}, \text{balance_score}\}$$

Flagging Rule:

$$\text{flagged} = 1(\text{bias_score} > 0.5)$$

References

- [Barocas et al., 2019] Fairness and Machine Learning
- [Hardt et al., 2016] Equality of Opportunity in Supervised Learning
- [Kusner et al., 2017] Counterfactual Fairness
- [Crenshaw, 1989] Demarginalizing the Intersection of Race and Sex

Prompt Injection Prevention

Overview

Multi-layered prevention system combining input sanitization, context isolation, and encoding detection.

Components

1. Input Sanitization

Method: Pattern removal while preserving semantic meaning

$$\text{sanitized}(x) = \text{remove_patterns}(x, P)$$

where **P** = set of injection patterns (system markers, separators, etc.)

Patterns Removed:

- System markers: <|system|>, [INST], </system>
- Separators: ---, ===, ***
- Role markers: user:, system:, assistant:

2. Context Isolation

Method: Strict separation of user input from system prompt

$$\text{isolated}(x, s) = (\text{clean}(x), s)$$

where:

- **x** = user input
- **s** = system prompt
- **clean(·)** = removal of system prompt markers from user input

3. Encoding Detection

Method: Detection of obfuscation techniques

$$\text{encoding_score} = \sum_{t \in T} w_t \cdot 1(\text{detect}_t(x))$$

where **T** = {URL encoding, hex encoding, Unicode tricks}

Weights: $w_{\text{URL}} = 0.3$, $w_{\text{hex}} = 0.3$, $w_{\text{unicode}} = 0.2$

4. Length Validation

Method: Detection of excessive length (potential injection)

$$\text{length_flag} = 1(\text{len}(x) > 10000)$$

Sanitization Pipeline

$$x' = \text{normalize}(\text{decode}(\text{remove_patterns}(x)))$$

where:

- **remove_patterns** = removes injection markers
- **decode** = decodes encoding tricks
- **normalize** = normalizes Unicode, whitespace

References

- [Greshake et al., 2023] Not What You've Signed Up For: Compromising Real-World LLM-Integrated Applications with Prompt Injection
- [Goodside, 2023] Prompt Injection Attacks Against GPT-3
- [OWASP, 2023] LLM Top 10: Prompt Injection

Dataset Testing Framework

Overview

Comprehensive evaluation framework for testing all safety components on standardized datasets.

Datasets

Toxicity Detection

1. Jigsaw Toxic Comment Classification (Kaggle)

- Size: ~159,000 comments
- Categories: Toxic, Severe Toxic, Obscene, Threat, Insult, Identity Hate
- Split: Train (80%), Validation (10%), Test (10%)

2. **Civil Comments** (Conversation AI)

- Size: ~2M comments
- Categories: Toxicity, Identity attacks, Insults, Threats
- Metadata: Demographics, identity groups

3. **Wikipedia Toxic Comments**

- Size: ~180,000 comments
- Categories: Toxic, Severe toxic, Obscene, Threat, Insult, Identity hate

Hallucination Detection

1. **FEVER** (Fact Extraction and VERification)

- Size: ~185,000 claims
- Labels: SUPPORTED, REFUTED, NOT ENOUGH INFO
- Use: Factuality checking

2. **XSum** (BBC News Summaries)

- Size: ~226,000 summaries
- Use: Consistency checking

3. **TruthfulQA**

- Size: ~800 questions
- Focus: Truthfulness and hallucinations

Jailbreak Detection

1. **Jailbreak Prompts Dataset** (Custom)

- Size: ~500 jailbreak examples
- Categories: Instruction override, Role manipulation, Adversarial

2. **HarmfulQA** (Safety Benchmarks)

- Size: ~1,000 harmful prompts
- Categories: Violence, Self-harm, Illegal activities

Bias Detection

1. **BOLD** (Bias in Open-Ended Language Generation Dataset)

- Size: ~23,000 prompts
- Demographics: Gender, Race, Religion, Profession

2. **StereoSet**

- Size: ~17,000 sentences
- Focus: Stereotype detection

3. **CrowS-Pairs**

- Size: ~1,500 sentence pairs
- Focus: Social bias across demographics

Evaluation Metrics

Classification Metrics

Accuracy:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

Precision:

$$\text{Precision} = \frac{TP}{TP + FP}$$

Recall:

$$\text{Recall} = \frac{TP}{TP + FN}$$

F1-Score:

$$F_1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Area Under ROC Curve (AUC-ROC):

$$\text{AUC} = \int_0^1 \text{TPR}(\text{FPR}^{-1}(t)) dt$$

Fairness Metrics

Demographic Parity:

$$\text{DP} = |P(\hat{Y} = 1|A = a) - P(\hat{Y} = 1|A = b)|$$

where **A** = protected attribute, **Ŷ** = prediction.

Equalized Odds:

$$\text{EO} = |P(\hat{Y} = 1|A = a, Y = y) - P(\hat{Y} = 1|A = b, Y = y)|$$

Calibration:

$$\text{Calibration} = \mathbb{E}[Y|\hat{P} = p] - p$$

Ideal: Calibration = 0 (perfect calibration)

Testing Framework Structure

```
tests/
└── datasets/
    ├── toxicity/
    │   ├── jigsaw_test.csv
    │   ├── civil_comments_test.csv
    │   └── wikipedia_test.csv
    ├── hallucination/
    │   ├── fever_test.jsonl
    │   ├── xsum_test.jsonl
    │   └── truthfulqa_test.jsonl
    ├── jailbreak/
    │   ├── jailbreak_prompts.jsonl
    │   └── harmfulqa_test.jsonl
    └── bias/
        ├── bold_test.jsonl
        ├── stereoset_test.jsonl
        └── crowdspairs_test.jsonl
├── test_toxicity.py
├── test_hallucination.py
├── test_jailbreak.py
├── test_bias.py
└── test_integration.py
```

Evaluation Scripts

Batch Evaluation:

```
# Pseudocode
for dataset in datasets:
    results = evaluate_model(model, dataset)
    metrics = calculate_metrics(results)
    save_results(dataset.name, metrics)
```

Cross-Dataset Evaluation:

$$\text{Generalization_Score} = \frac{1}{n} \sum_{i=1}^n F_1(\text{model}, D_i)$$

where D_i = dataset i , n = number of datasets.

References

- [Borkan et al., 2019] Nuanced Metrics for Measuring Unintended Bias with Real Data for Text Classification
- [Thorne et al., 2018] FEVER: a Large-scale Dataset for Fact Extraction and VERification
- [Nadeem et al., 2021] StereoSet: Measuring Stereotypical Bias in Pretrained Language Models
- [Nangia et al., 2020] CrowS-Pairs: A Challenge Dataset for Measuring Social Bias in Masked Language Models

Performance Metrics

Latency Targets

- **Toxicity Detection:** <50ms (GPU), <200ms (CPU)
- **Hallucination Detection:** <300ms (multi-component)
- **Jailbreak Detection:** <100ms
- **Bias Detection:** <500ms (with demographic testing)

Throughput

- **Batch Processing:** >100 requests/second
- **Concurrent Requests:** >50 parallel requests
- **GPU Utilization:** >80% (under load)

Accuracy Targets

- **Toxicity F1:** >0.92
- **Hallucination AUC:** >0.85
- **Jailbreak Recall:** >0.95 (high recall for security)
- **Bias Detection:** >0.80 F1

Resource Requirements

- **GPU:** NVIDIA GPU with 8GB+ VRAM (recommended)
- **CPU:** 4+ cores, 16GB+ RAM (minimum)
- **Storage:** 10GB+ for models and datasets

Implementation Timeline

Phase 1: Core Detection (Weeks 1-4)

- Custom toxicity detector
- Basic hallucination detection
- Advanced jailbreak detection
- Enhanced bias detection

Phase 2: Testing Framework (Weeks 5-6)

- Dataset preparation
- Evaluation scripts
- Benchmarking suite
- Integration tests

Phase 3: Optimization (Weeks 7-8)

- Model quantization
- Batch processing optimization
- Caching layer
- Performance tuning

Phase 4: Production Readiness (Weeks 9-10)

- Monitoring and logging
- Error handling
- Documentation
- Deployment scripts

References

Toxicity Detection

1. Devlin, J., et al. (2019). "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding." *NAACL-HLT*.
2. Liu, Y., et al. (2019). "RoBERTa: A Robustly Optimized BERT Pretraining Approach." *arXiv preprint arXiv:1907.11692*.
3. Unitary AI. (2021). "Toxic-BERT: Multi-label Toxicity Classification." *HuggingFace Model Card*.

Hallucination Detection

4. Maynez, J., et al. (2020). "On Faithfulness and Factuality in Abstractive Summarization." *ACL*.
5. Kryściński, W., et al. (2020). "Evaluating the Factual Consistency of Abstractive Text Summarization." *EMNLP*.
6. Reimers, N., & Gurevych, I. (2019). "Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks." *EMNLP*.
7. He, P., et al. (2021). "DeBERTa: Decoding-enhanced BERT with Disentangled Attention." *ICLR*.

Jailbreak Detection

8. Perez, E., et al. (2022). "Red Teaming Language Models to Reduce Harms." *arXiv preprint arXiv:2209.07858*.
9. Zou, A., et al. (2023). "Universal and Transferable Adversarial Attacks on Aligned Language Models." *arXiv preprint arXiv:2307.15043*.
10. Wei, A., et al. (2023). "Jailbroken: How Does LLM Safety Training Fail?" *arXiv preprint arXiv:2307.02483*.

Bias Detection

11. Barocas, S., et al. (2019). "Fairness and Machine Learning." fairmlbook.org.
12. Hardt, M., et al. (2016). "Equality of Opportunity in Supervised Learning." *NeurIPS*.
13. Kusner, M. J., et al. (2017). "Counterfactual Fairness." *NeurIPS*.
14. Crenshaw, K. (1989). "Demarginalizing the Intersection of Race and Sex: A Black Feminist Critique of Antidiscrimination Doctrine." *University of Chicago Legal Forum*.

Prompt Injection

15. Greshake, K., et al. (2023). "Not What You've Signed Up For: Compromising Real-World LLM-Integrated Applications with Prompt Injection." *arXiv preprint arXiv:2303.18103*.
16. Goodside, R. (2023). "Prompt Injection Attacks Against GPT-3." *Personal Blog*.
17. OWASP. (2023). "LLM Top 10: Prompt Injection." *OWASP Foundation*.

Dataset Testing

18. Borkan, D., et al. (2019). "Nuanced Metrics for Measuring Unintended Bias with Real Data for Text Classification." *ACL*.
19. Thorne, J., et al. (2018). "FEVER: a Large-scale Dataset for Fact Extraction and VERification." *NAACL*.
20. Nadeem, M., et al. (2021). "StereoSet: Measuring Stereotypical Bias in Pretrained Language Models." *ACL*.
21. Nangia, N., et al. (2020). "CrowS-Pairs: A Challenge Dataset for Measuring Social Bias in Masked Language Models." *EMNLP*.

Appendix: Mathematical Notation

- **x, y, z:** Vectors or scalars (context-dependent)
- **X, Y, Z:** Matrices or sets
- **p:** Probability or prompt (context-dependent)
- **P(·):** Probability distribution
- **cos(·, ·):** Cosine similarity
- **embed(·):** Embedding function
- **σ(·):** Sigmoid function
- **E[·]:** Expectation
- **Var(·):** Variance
- **TP, TN, FP, FN:** True/False Positives/Negatives
- **τ:** Threshold value
- **w_i:** Weight for component **i**
- **s_i:** Score for component **i**

Document Status: Living Document - Updated as implementation progresses

Contact: av.rajpurkar@gmail.com

Repository: <https://github.com/aryan4codes/Equitas>